

Retail Data Analysis

Code Logic

Logic for Python Script 'spark-streaming.py'

Setting up the system dependencies for Cloudera distribution by importing necessary libraries, modules and the path variables

```
import os
import sys

os.environ["PYSPARK_PYTHON"] = "/opt/cloudera/parcels/Anaconda/bin/python"
os.environ["JAVA_HOME"] = "/usr/java/jdk1.8.0_232-cloudera/jre"
os.environ["SPARK_HOME"]="/opt/cloudera/parcels/SPARK2-2.3.0.cloudera2-1.cdh5.13.3.p0.316101/lib/spark2/"
os.environ["PYLIB"] = os.environ["SPARK_HOME"] + "/python/lib"
sys.path.insert(0, os.environ["PYLIB"] + "/py4j-0.10.6-src.zip")
sys.path.insert(0, os.environ["PYLIB"] + "/pyspark.zip")

from pyspark.sql import SparkSession
from pyspark.sql.functions import *
from pyspark.sql.types import *
```

Writing the Python functions, which contain the logic for the UDFs

1. Total Cost UDF - To calculate the total income from every invoice I needed to calculate the income from sale of each product, so I multiplied the unit price of the product with the quantity of the product purchased. The sum of this cost across the products in that invoice gives me the total cost of the order. I also made sure that if the transaction is a return transaction, the total cost is negative.

```
def find_total_order_cost(items, trn_type):
    if items is not None:
        total_cost = 0
        item_price = 0
        for item in items:
            item_price = (item['quantity'] * item['unit_price'])
            total_cost = total_cost + item_price
            item_price = 0

        if trn_type == "RETURN":
            return total_cost * -1
        else:
            return total_cost
```

2. Total Items UDF - To calculate the number of products in every invoice I added the quantity ordered of each product in that invoice

```
def find_total_item_count(items):
    if items is not None:
        total_count = 0
        for item in items:
            total_count = total_count + item['quantity']
        return total_count
```

3. Is Order UDF - To determine if invoice is for an order or not I used an if-else statement

```
def flag_isOrder(trn_type):  
    if trn_type == "ORDER":  
        return(1)  
    else:  
        return(0)
```

4. Is Return UDF - To determine if invoice is for a return or not I used an if-else statement

```
def flag_isReturn(trn_type):  
    if trn_type == "RETURN":  
        return(1)  
    else:  
        return(0)
```

Initialising the Spark session and setting the log level to error as a good practice

```
spark = SparkSession \  
    .builder \  
    .appName("spark-streaming") \  
    .getOrCreate()  
spark.sparkContext.setLogLevel('ERROR')
```

Reading input data from Kafka mentioning the details of the Kafka broker, such as bootstrap server, port and topic name

```
orderRawData = spark.readStream \  
    .format("kafka") \  
    .option("kafka.bootstrap.servers", "18.211.252.152:9092") \  
    .option("startingOffsets", "earliest") \  
    .option("failOnDataLoss", "false") \  
    .option("subscribe", "real-time-project") \  
    .load()
```

Defining JSON schema of each order, using appropriate datatypes and StructField in the case of the item attributes

```
jsonSchema = StructType() \  
    .add("invoice_no", LongType()) \  
    .add("country", StringType()) \  
    .add("timestamp", TimestampType()) \  
    .add("type", StringType()) \  
    .add("items", ArrayType(StructType([  
        StructField("SKU", StringType()),  
        StructField("title", StringType()),  
        StructField("unit_price", FloatType()),  
        StructField("quantity", IntegerType()),  
    ])))
```

Reading the raw JSON data from Kafka as 'order stream' by casting it to string and storing it into the alias 'data'

```
orderStream = orderRawData.select(from_json(col("value").cast("string"),  
jsonSchema).alias("data")).select("data.*")
```

Defining the UDFs by Converting the Python functions I defined earlier, and assigning the appropriate return datatype

```
sum_total_order_cost = udf(find_total_order_cost, FloatType())
sum_total_item_count = udf(find_total_item_count, IntegerType())
sum_isOrder = udf(flag_isOrder, IntegerType())
sum_isReturn = udf(flag_isReturn, IntegerType())
```

Calculating the additional columns according to the required input values

```
expandedOrderStream = orderStream \
    .withColumn("total_cost", sum_total_order_cost(orderStream.items,
orderStream.type)) \
    .withColumn("total_items", sum_total_item_count(orderStream.items)) \
    .withColumn("is_order", sum_isOrder(orderStream.type)) \
    .withColumn("is_return", sum_isReturn(orderStream.type))
```

Writing the summarised input values to console, using 'append' output method and applying truncate as false and setting the processing time to 1 minute

```
extendedOrderQuery = expandedOrderStream \
    .select("invoice_no", "country", "timestamp", "total_cost", "total_items",
"is_order", "is_return") \
    .writeStream \
    .outputMode("append") \
    .format("console") \
    .option("truncate", "false") \
    .trigger(processingTime = "1 minute") \
    .start()
```

Calculating time-based KPIs (Total sale volume, OPM, Rate of return, Average transaction size) having tumbling window of one minute and watermark of one minute.

```
aggStreamByTime = expandedOrderStream \
    .withWatermark("timestamp", "1 minute") \
    .groupBy(window("timestamp", "1 minute", "1 minute")) \
    .agg(sum("total_cost").alias("total_sale_volume"),
        count("invoice_no").alias("OPM"),
        avg("is_return").alias("rate_of_return"),
        avg("total_cost").alias("average_transaction_size")
    ) \
    .select("window", "OPM", "total_sale_volume", "average_transaction_size",
"rate_of_return")
```

Writing the time-based KPIs data to HDFS - HDFS into JSON files for each one-minute window, using 'append' output mode, setting truncate as false, and specifying the HDFS output path for both the KPI files and for their checkpoints. Ten 1-minute window batches were taken.

```
queryByTime = aggStreamByTime.writeStream \
    .format("json") \
    .outputMode("append") \
    .option("truncate", "false") \
    .option("path", "/user/ec2-user/time_kpi") \
    .option("checkpointLocation", "/user/ec2-user/time_kpi_checkpoints") \
    .trigger(processingTime="1 minute") \
    .start()
```

Calculating time-and-country-based KPIs (Total sale volume, OPM, Rate of return) having tumbling window of one minute and watermark of one minute. Here I grouped by window and country both.

```
aggStreamByCountry = expandedOrderStream \  
    .withWatermark("timestamp", "1 minute") \  
    .groupBy(window("timestamp", "1 minute", "1 minute"), "country") \  
    .agg(sum("total_cost").alias("total_sale_volume"),  
        count("invoice_no").alias("OPM"),  
        avg("is_return").alias("rate_of_return")) \  
    .select("window", "country", "OPM", "total_sale_volume", "rate_of_return")
```

Writing the the time-and-country-based KPIs data to HDFS into JSON files for each one-minute window, using 'append' output mode, setting truncate as false, and specifying the HDFS output path for both the KPI files and for their checkpoints. Ten 1-minute window batches were taken.

```
queryByCountry = aggStreamByCountry.writeStream \  
    .format("json") \  
    .outputMode("append") \  
    .option("truncate","false") \  
    .option("path", "/user/ec2-user/country_kpi") \  
    .option("checkpointLocation", "/user/ec2-user/country_kpi_checkpoints") \  
    .trigger(processingTime="1 minute") \  
    .start()
```

Indicating Spark to await termination

```
extendedOrderQuery.awaitTermination()  
queryByCountry.awaitTermination()  
queryByTime.awaitTermination()
```

Console Commands

I started by logging into the ec2 instance as 'ec2-user'

Next, I downloaded the Spark-SQL-Kafka jar file. This jar is used to run the Spark Streaming-Kafka codes

```
wget https://ds-spark-sql-kafka-jar.s3.amazonaws.com/spark-sql-kafka-0-10_2.11-  
2.3.0.jar
```

Next, I created the 'spark-streaming.py' file having the code discussed above

```
vi spark-streaming.py
```

Next, I set the Kafka Version using the following command

```
export SPARK_KAFKA_VERSION=0.10
```

Finally, I ran the spark2-submit command, specifying the jar and python file

```
spark2-submit --jars spark-sql-kafka-0-10_2.11-2.3.0.jar spark-streaming.py
```

Example table - Final Summarised Input Values

Batch: 0

invoice_no	country	timestamp	total_cost	total_items	is_order	is_return
154132552443909	United Kingdom	2022-12-08 08:09:18	93.52	48	1	0
154132552443910	United Kingdom	2022-12-08 08:09:19	143.73	103	1	0
154132552443911	Belgium	2022-12-08 08:09:45	171.09001	195	1	0
154132552443912	United Kingdom	2022-12-08 08:09:56	-14.22	6	0	1
154132552443913	United Kingdom	2022-12-08 08:09:58	67.689995	13	1	0
154132552443914	United Kingdom	2022-12-08 08:09:59	42.63	35	1	0
154132552443915	United Kingdom	2022-12-08 08:10:00	0.42	1	1	0
154132552443916	United Kingdom	2022-12-08 08:10:05	25.93	43	1	0
154132552443917	United Kingdom	2022-12-08 08:10:08	17.85	4	1	0
154132552443918	United Kingdom	2022-12-08 08:10:12	-19.289999	6	0	1
154132552443919	United Kingdom	2022-12-08 08:10:15	749.83	405	1	0
154132552443920	United Kingdom	2022-12-08 08:10:21	28.34	10	1	0
154132552443921	United Kingdom	2022-12-08 08:10:29	133.84	67	1	0
154132552443922	United Kingdom	2022-12-08 08:10:33	17.15	3	1	0
154132552443923	United Kingdom	2022-12-08 08:10:34	43.1	28	1	0
154132552443924	United Kingdom	2022-12-08 08:10:37	21.04	8	1	0
154132552443925	United Kingdom	2022-12-08 08:10:38	70.36	43	1	0
154132552443926	United Kingdom	2022-12-08 08:10:55	19.93	5	1	0
154132552443927	United Kingdom	2022-12-08 08:11:06	20.0	16	1	0
154132552443928	United Kingdom	2022-12-08 08:11:07	100.29	75	1	0

only showing top 20 rows

Batch: 1

invoice_no	country	timestamp	total_cost	total_items	is_order	is_return
154132552467961	United Kingdom	2022-12-09 23:57:15	227.87999	26	1	0
154132552467962	United Kingdom	2022-12-09 23:57:15	-18.82	20	0	1
154132552467963	United Kingdom	2022-12-09 23:57:20	126.89	37	1	0
154132552467964	United Kingdom	2022-12-09 23:57:24	19.56	12	1	0
154132552467965	United Kingdom	2022-12-09 23:57:25	37.86	13	1	0
154132552467966	United Kingdom	2022-12-09 23:57:26	47.88	11	1	0
154132552467967	United Kingdom	2022-12-09 23:57:35	10.5	25	1	0
154132552467968	Germany	2022-12-09 23:57:41	18.15	11	1	0
154132552467969	United Kingdom	2022-12-09 23:57:44	18.4	4	1	0
154132552467970	United Kingdom	2022-12-09 23:57:45	56.690002	53	1	0
154132552467971	United Kingdom	2022-12-09 23:57:46	16.06	28	1	0
154132552467972	United Kingdom	2022-12-09 23:57:47	141.6	48	1	0
154132552467973	United Kingdom	2022-12-09 23:57:47	31.810001	3	1	0
154132552467974	United Kingdom	2022-12-09 23:57:59	4.25	1	1	0

I checked HDFS to make sure the KPI files were present

```
hadoop fs -ls /user/ec2-user
drwxr-xr-x - hadoop hadoop 0 2022-04-12 19:05 /user/hadoop/.sparkStaging
drwxr-xr-x - hadoop hadoop 0 2022-04-12 19:05 /user/hadoop/country_kpi
drwxr-xr-x - hadoop hadoop 0 2022-04-12 18:46 /user/hadoop/country_kpi_checkpoints
drwxr-xr-x - hadoop hadoop 0 2022-04-12 19:05 /user/hadoop/time_kpi
drwxr-xr-x - hadoop hadoop 0 2022-04-12 18:46 /user/hadoop/time_kpi_checkpoints
```

I also checked the folders to see the JSON files

```
hadoop fs -ls /user/ec2-user/time_kpi/
```



```
Found 240 items
drwxr-xr-x - hadoop hadoop 0 2022-04-12 19:05 /user/hadoop/time_kpi/ spark metadata
-rw-r--r-- 1 hadoop hadoop 6560 2022-04-12 18:46 /user/hadoop/time_kpi/part-00000-149396a6-1f83-42ce-a5e9-e2c012d5110a-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:50 /user/hadoop/time_kpi/part-00000-34bb2c5a-a0e2-4254-b80d-0fb248c41446-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:47 /user/hadoop/time_kpi/part-00000-3cb0d6c2-a8c8-495e-9c25-301331692bc4-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:54 /user/hadoop/time_kpi/part-00000-4849c1a7-9755-4e46-8544-17a91568af10-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:56 /user/hadoop/time_kpi/part-00000-563349d2-1d2c-4355-884d-dd9c94ef86a5-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:58 /user/hadoop/time_kpi/part-00000-5c3aae37-e1b6-4e94-bc6d-4a66348578ba-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:46 /user/hadoop/time_kpi/part-00000-60967d50-29ee-4ce8-ab8b-3fddc38adc6d-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 19:02 /user/hadoop/time_kpi/part-00000-61691cfb-e449-41ff-8c92-dd3de2b7da9f-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 19:03 /user/hadoop/time_kpi/part-00000-80b32683-d514-4892-b7e9-dd302060fd6e-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 19:05 /user/hadoop/time_kpi/part-00000-821872c2-230b-4457-afd2-d2a5932c2cec-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:48 /user/hadoop/time_kpi/part-00000-87cacd49-8f5e-4639-9459-69f338302be1-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 19:00 /user/hadoop/time_kpi/part-00000-8f47bd4c-f0eb-4289-af1b-f960fc575af3-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:51 /user/hadoop/time_kpi/part-00000-9fc309ca-cf5e-45bc-b7b4-1ce716db7988-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 19:01 /user/hadoop/time_kpi/part-00000-a0633b82-8123-4784-a9bc-6f2d1a2e99c6-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:59 /user/hadoop/time_kpi/part-00000-af3430f7-c9eb-43a5-955f-31dba49dc9fc-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:55 /user/hadoop/time_kpi/part-00000-c759baba-3248-4a37-903d-2eb22f5a5b5d-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:57 /user/hadoop/time_kpi/part-00000-e7068283-9a0c-4747-a0d8-c5f6df92c7a2-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:49 /user/hadoop/time_kpi/part-00000-e91cb7c2-1784-4fea-9ca8-508b5083bda0-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 19:04 /user/hadoop/time_kpi/part-00000-ea31e354-7a49-49f3-ade5-483ed10df024-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:53 /user/hadoop/time_kpi/part-00000-fadcd7a-f764-4a01-a63c-77136e41547d-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:52 /user/hadoop/time_kpi/part-00000-fb48ff4-88ee-4716-a63c-d8d1b33e0295-c000.json
-rw-r--r-- 1 hadoop hadoop 8250 2022-04-12 18:46 /user/hadoop/time_kpi/part-00001-26bb2da2-0fd8-41ae-a928-224fca0a765f-c000.json
-rw-r--r-- 1 hadoop hadoop 7181 2022-04-12 18:46 /user/hadoop/time_kpi/part-00002-8c2607eb-5a5b-4bc2-9952-88af43001942-c000.json
-rw-r--r-- 1 hadoop hadoop 7684 2022-04-12 18:46 /user/hadoop/time_kpi/part-00003-8669f8b5-907d-4c1c-8930-1870145a2b71-c000.json
-rw-r--r-- 1 hadoop hadoop 6494 2022-04-12 18:46 /user/hadoop/time_kpi/part-00004-1a59c3a1-c1e4-4f2b-96f0-4f3af3c76dal-c000.json
-rw-r--r-- 1 hadoop hadoop 7692 2022-04-12 18:46 /user/hadoop/time_kpi/part-00005-2833c712-4927-44ca-a913-63ba5e0621e4-c000.json
-rw-r--r-- 1 hadoop hadoop 7815 2022-04-12 18:46 /user/hadoop/time_kpi/part-00006-3f7cb50a-8595-4efe-9804-50eca72e4f3d-c000.json
-rw-r--r-- 1 hadoop hadoop 8313 2022-04-12 18:46 /user/hadoop/time_kpi/part-00007-b39f55ab-82a9-48c9-814c-321d59d72eaa-c000.json
-rw-r--r-- 1 hadoop hadoop 8209 2022-04-12 18:46 /user/hadoop/time_kpi/part-00008-1d432b84-657f-4667-9326-6126a46d358c-c000.json
-rw-r--r-- 1 hadoop hadoop 7633 2022-04-12 18:46 /user/hadoop/time_kpi/part-00009-f0a537a8-3220-4c74-87a3-f58bee6f1e0d-c000.json
-rw-r--r-- 1 hadoop hadoop 8185 2022-04-12 18:46 /user/hadoop/time_kpi/part-00010-6ada4f0f-0c6b-4b9a-acb6-4c354678d161-c000.json
```

hadoop fs -ls /user/ec2-user/country_kpi/

```
[hadoop@ip-172-31-40-71 ~]$ hadoop fs -ls /user/hadoop/country_kpi/
Found 254 items
drwxr-xr-x - hadoop hadoop 0 2022-04-12 19:05 /user/hadoop/country_kpi/ spark metadata
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:46 /user/hadoop/country_kpi/part-00000-2266d962-5f3c-43a9-9ddb-64609c5f58e2-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 19:05 /user/hadoop/country_kpi/part-00000-22aa280a-8e16-4d65-a7b9-f4371df93086-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:51 /user/hadoop/country_kpi/part-00000-2c6fe6e7-a975-4ab8-8eb5-71a883f01e2f-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:58 /user/hadoop/country_kpi/part-00000-2ee4f36f-3717-4433-9fdb-8a81f9d2429b-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:55 /user/hadoop/country_kpi/part-00000-435c9bd3-abcb-4982-a684-e9e50c4be98e-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:52 /user/hadoop/country_kpi/part-00000-4dd52f3c-71ff-4a30-8f89-d133a2f7ef5e-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:57 /user/hadoop/country_kpi/part-00000-5087bd7f-7ae1-40a2-af28-4145827ffa24-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 19:02 /user/hadoop/country_kpi/part-00000-6f389c54-0790-4b67-acca-4d4983c5a583-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:56 /user/hadoop/country_kpi/part-00000-8363420d-fbe4-4a8e-bf6c-b596d68ea5c3-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:48 /user/hadoop/country_kpi/part-00000-92714469-f16e-46c9-b72f-fa0b1048899ac-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 19:04 /user/hadoop/country_kpi/part-00000-962e6455-3f6e-4b59-8357-0ee2fe1899d4-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 19:00 /user/hadoop/country_kpi/part-00000-98fe621c-958e-4aa4-88aa-clc31d3a5b0d-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:54 /user/hadoop/country_kpi/part-00000-a8f7bcf6-5708-430b-9af4-f9448d00450f-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:59 /user/hadoop/country_kpi/part-00000-b926cf9d-c496-4af5-8cf3-878c6961c191-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:50 /user/hadoop/country_kpi/part-00000-bad9a499-2c67-4ea7-b9d2-d80f4de427d5-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:53 /user/hadoop/country_kpi/part-00000-d348f141-d19f-42fe-a01a-6ce71f2006e8-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:49 /user/hadoop/country_kpi/part-00000-d925a3ca-4be1-47d4-a58c-02c3e179a7a2-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 18:47 /user/hadoop/country_kpi/part-00000-e781261e-4be7-40c3-bf5b-c734a30f6d1a-c000.json
-rw-r--r-- 1 hadoop hadoop 12794 2022-04-12 18:46 /user/hadoop/country_kpi/part-00000-f767ddfc-e47b-4520-ab91-c9fba4627434-c000.json
-rw-r--r-- 1 hadoop hadoop 0 2022-04-12 19:03 /user/hadoop/country_kpi/part-00000-fa74d646-e10b-48ce-8819-16fdd50a1578-c000.json
```

And used 'cat' command to take a look at the data

hadoop fs -cat /user/ec2-user/time_kpi/part*

```
[ec2-user@ip-10-0-0-71 ~]$ hadoop fs -cat /user/ec2-user/time_kpi/part*
{"window":{"start":"2021-10-25T11:15:00.000Z","end":"2021-10-25T11:16:00.000Z"},"OPM":9,"total_sale_volume":245.0800018310547,"average_transaction_size":27.2311131456163,"rate_of_return":0.111111111111111}
{"window":{"start":"2021-10-25T11:24:00.000Z","end":"2021-10-25T11:25:00.000Z"},"OPM":12,"total_sale_volume":1079.4999951422215,"average_transaction_size":89.9583322851846,"rate_of_return":0.0833333333333333}
{"window":{"start":"2021-10-25T11:19:00.000Z","end":"2021-10-25T11:20:00.000Z"},"OPM":11,"total_sale_volume":428.7600100636482,"average_transaction_size":39.87818273305893,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:13:00.000Z","end":"2021-10-25T11:14:00.000Z"},"OPM":9,"total_sale_volume":459.8600025177002,"average_transaction_size":51.09555585350002,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:13:00.000Z","end":"2021-10-25T11:14:00.000Z"},"OPM":9,"total_sale_volume":459.8600025177002,"average_transaction_size":51.09555585350002,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:14:00.000Z","end":"2021-10-25T11:15:00.000Z"},"OPM":14,"total_sale_volume":483.4499945640564,"average_transaction_size":34.532142468861174,"rate_of_return":0.07142857142857142}
{"window":{"start":"2021-10-25T11:11:00.000Z","end":"2021-10-25T11:12:00.000Z"},"OPM":5,"total_sale_volume":396.11000061035156,"average_transaction_size":79.22200012207031,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:20:00.000Z","end":"2021-10-25T11:21:00.000Z"},"OPM":5,"total_sale_volume":258.56000328063965,"average_transaction_size":51.71200065612793,"rate_of_return":0.2}
{"window":{"start":"2021-10-25T11:17:00.000Z","end":"2021-10-25T11:18:00.000Z"},"OPM":7,"total_sale_volume":1067.4699907302856,"average_transaction_size":152.49571296146937,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:12:00.000Z","end":"2021-10-25T11:13:00.000Z"},"OPM":13,"total_sale_volume":468.60000121593475,"average_transaction_size":36.04615393968729,"rate_of_return":0.07692307692307693}
{"window":{"start":"2021-10-25T11:18:00.000Z","end":"2021-10-25T11:19:00.000Z"},"OPM":7,"total_sale_volume":1583.7199920415878,"average_transaction_size":226.24571314879827,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:21:00.000Z","end":"2021-10-25T11:22:00.000Z"},"OPM":14,"total_sale_volume":668.9900054335594,"average_transaction_size":47.78500038811385,"rate_of_return":0.14285714285714285}
{"window":{"start":"2021-10-25T11:16:00.000Z","end":"2021-10-25T11:17:00.000Z"},"OPM":4,"total_sale_volume":78.17999869585037,"average_transaction_size":19.544999673962593,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:25:00.000Z","end":"2021-10-25T11:26:00.000Z"},"OPM":15,"total_sale_volume":542.6700085401535,"average_transaction_size":36.17800056934357,"rate_of_return":0.1333333333333333}
{"window":{"start":"2021-10-25T11:23:00.000Z","end":"2021-10-25T11:24:00.000Z"},"OPM":15,"total_sale_volume":469.79999724030495,"average_transaction_size":31.31999981602033,"rate_of_return":0.06666666666666667}
{"window":{"start":"2021-10-25T11:22:00.000Z","end":"2021-10-25T11:23:00.000Z"},"OPM":12,"total_sale_volume":335.83000135421753,"average_transaction_size":27.985833446184795,"rate_of_return":0.0}
```

hadoop fs -cat /user/ec2-user/country_kpi/part*

```
[ec2-user@ip-10-0-0-71 ~]$ hadoop fs -cat /user/ec2-user/country_kpi/part*
{"window":{"start":"2021-10-25T11:20:00.000Z","end":"2021-10-25T11:21:00.000Z"},"country":"Unspecified","OPM":1,"total_sale_volume":-93.43000030517578,"rate_of_return":1.0}
{"window":{"start":"2021-10-25T11:26:00.000Z","end":"2021-10-25T11:27:00.000Z"},"country":"EIRE","OPM":1,"total_sale_volume":29.15999984741211,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:24:00.000Z","end":"2021-10-25T11:25:00.000Z"},"country":"Denmark","OPM":1,"total_sale_volume":157.77999877929688,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:21:00.000Z","end":"2021-10-25T11:22:00.000Z"},"country":"Channel Islands","OPM":1,"total_sale_volume":23.1299991607666,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:19:00.000Z","end":"2021-10-25T11:20:00.000Z"},"country":"United Kingdom","OPM":11,"total_sale_volume":428.7600100636482,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:23:00.000Z","end":"2021-10-25T11:24:00.000Z"},"country":"United Kingdom","OPM":15,"total_sale_volume":469.79999724030495,"rate_of_return":0.06666666666666667}
{"window":{"start":"2021-10-25T11:25:00.000Z","end":"2021-10-25T11:26:00.000Z"},"country":"France","OPM":2,"total_sale_volume":30.779999136924744,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:22:00.000Z","end":"2021-10-25T11:23:00.000Z"},"country":"Spain","OPM":1,"total_sale_volume":47.06999969482422,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:26:00.000Z","end":"2021-10-25T11:27:00.000Z"},"country":"United Kingdom","OPM":4,"total_sale_volume":228.52000045776367,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:17:00.000Z","end":"2021-10-25T11:18:00.000Z"},"country":"Norway","OPM":1,"total_sale_volume":11.899999618530273,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:12:00.000Z","end":"2021-10-25T11:13:00.000Z"},"country":"EIRE","OPM":1,"total_sale_volume":4.440000057220459,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:14:00.000Z","end":"2021-10-25T11:15:00.000Z"},"country":"United Kingdom","OPM":14,"total_sale_volume":483.4499945640564,"rate_of_return":0.07142857142857142}
{"window":{"start":"2021-10-25T11:13:00.000Z","end":"2021-10-25T11:14:00.000Z"},"country":"Germany","OPM":1,"total_sale_volume":38.59000015258789,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:12:00.000Z","end":"2021-10-25T11:13:00.000Z"},"country":"United Kingdom","OPM":12,"total_sale_volume":464.160011587143,"rate_of_return":0.08333333333333333}
{"window":{"start":"2021-10-25T11:21:00.000Z","end":"2021-10-25T11:22:00.000Z"},"country":"United Kingdom","OPM":13,"total_sale_volume":645.860062727928,"rate_of_return":0.15384615384615385}
{"window":{"start":"2021-10-25T11:18:00.000Z","end":"2021-10-25T11:19:00.000Z"},"country":"United Kingdom","OPM":7,"total_sale_volume":1583.7199920415878,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:22:00.000Z","end":"2021-10-25T11:23:00.000Z"},"country":"United Kingdom","OPM":11,"total_sale_volume":288.760016593933,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:13:00.000Z","end":"2021-10-25T11:14:00.000Z"},"country":"United Kingdom","OPM":7,"total_sale_volume":349.870000839234,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:16:00.000Z","end":"2021-10-25T11:17:00.000Z"},"country":"United Kingdom","OPM":4,"total_sale_volume":78.17999869585037,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:17:00.000Z","end":"2021-10-25T11:18:00.000Z"},"country":"United Kingdom","OPM":6,"total_sale_volume":1055.569911117554,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:24:00.000Z","end":"2021-10-25T11:25:00.000Z"},"country":"United Kingdom","OPM":10,"total_sale_volume":930.219963629246,"rate_of_return":0.0}
```

Transfer of files from CDH Instance on AWS to my system, using WinSCP

First, I needed to transfer the JSON files from HDFS into the the EC2 system

I created directories for time-based and then time-and-country-based KPIs as ec2-user. Using the 'get' command I copied the contents of the output folders into the EC2 system.

```
mkdir timebased-KPI
hadoop fs -get /user/ec2-user/time_kpi /home/ec2-user/timebased-KPI
```



```
mkdir country-and-timebased-KPI  
hadoop fs -get /user/ec2-user/country_kpi /home/ec2-user/country-and-timebased-KPI
```

Thereafter I used WinSCP to establish a connection between the EC2 instance and my local file system to transfer all the required files into my system.